

KFold Cross Validation Python Tutorial

```
In [83]: from sklearn.linear_model import LogisticRegression
        from sklearn.svm import SVC
        from sklearn.ensemble import RandomForestClassifier
        import numpy as np
        from sklearn.datasets import load_digits
        import matplotlib.pyplot as plt
        digits = load_digits()
```

```
In [84]: from sklearn.model_selection import train_test_split
        X_train, X_test, y_train, y_test = train_test_split(digits.data, digits.
        target, test_size=0.3)
```

Logistic Regression

```
In [85]: lr = LogisticRegression(solver='liblinear', multi_class='ovr')
        lr.fit(X_train, y_train)
        lr.score(X_test, y_test)
```

```
Out[85]: 0.9666666666666667
```

SVM

```
In [86]: svm = SVC(gamma='auto')
        svm.fit(X_train, y_train)
        svm.score(X_test, y_test)
```

```
Out[86]: 0.3685185185185185
```

Random Forest

```
In [87]: rf = RandomForestClassifier(n_estimators=40)
rf.fit(X_train, y_train)
rf.score(X_test, y_test)
```

```
Out[87]: 0.9703703703703703
```

KFold cross validation

Basic example

```
In [88]: from sklearn.model_selection import KFold
kf = KFold(n_splits = 3)
kf
```

```
Out[88]: KFold(n_splits=3, random_state=None, shuffle=False)
```

```
In [89]: for train_index, test_index in kf.split([1,2,3,4,5,6,7,8,9]):
print(train_index,test_index)
```

```
[3 4 5 6 7 8] [0 1 2]
[0 1 2 6 7 8] [3 4 5]
[0 1 2 3 4 5] [6 7 8]
```

Use KFold for our digits example

```
In [90]: def get_score(model,X_train,X_test,y_train,y_test):
model.fit(X_train,y_train)
return model.score(X_test,y_test)
```

```
In [91]: from sklearn.model_selection import StratifiedKFold
folds = StratifiedKFold(n_splits=3)

scores_logistic = []
scores_svm = []
scores_rf = []
```

```

for train_index, test_index in folds.split(digits.data,digits.target):
    X_train, X_test, y_train, y_test = digits.data[train_index], digits
    .data[test_index], \
                                digits.target[train_index], digi
    ts.target[test_index]
    scores_logistic.append(get_score(LogisticRegression(solver='libline
ar',multi_class='ovr'), X_train, X_test, y_train, y_test))
    scores_svm.append(get_score(SVC(gamma='auto'), X_train, X_test, y_t
rain, y_test))
    scores_rf.append(get_score(RandomForestClassifier(n_estimators=40),
X_train, X_test, y_train, y_test))

```

In [92]: scores_logistic

Out[92]: [0.8948247078464107, 0.9532554257095158, 0.9098497495826378]

In [93]: scores_svm

Out[93]: [0.3806343906510851, 0.41068447412353926, 0.5125208681135225]

In [94]: scores_rf

Out[94]: [0.9248747913188647, 0.9465776293823038, 0.9248747913188647]

cross_val_score function

In [95]: `from sklearn.model_selection import cross_val_score`

Logistic regression model performance using cross_val_score

In [97]: `cross_val_score(LogisticRegression(solver='liblinear',multi_class='ovr'), digits.data, digits.target,cv=3)`

Out[97]: array([0.89482471, 0.95325543, 0.90984975])

svm model performance using cross_val_score

```
In [98]: cross_val_score(SVC(gamma='auto'), digits.data, digits.target, cv=3)
```

```
Out[98]: array([0.38063439, 0.41068447, 0.51252087])
```

random forest performance using cross_val_score

```
In [99]: cross_val_score(RandomForestClassifier(n_estimators=40), digits.data, digits.target, cv=3)
```

```
Out[99]: array([0.92988314, 0.93656093, 0.92320534])
```

cross_val_score uses stratified kfold by default

Parameter tuning using k fold cross validation

```
In [100]: scores1 = cross_val_score(RandomForestClassifier(n_estimators=5), digits.data, digits.target, cv=10)
          np.average(scores1)
```

```
Out[100]: 0.892029795158287
```

```
In [101]: scores2 = cross_val_score(RandomForestClassifier(n_estimators=20), digits.data, digits.target, cv=10)
          np.average(scores2)
```

```
Out[101]: 0.9332091868404717
```

```
In [102]: scores3 = cross_val_score(RandomForestClassifier(n_estimators=30), digits.data, digits.target, cv=10)
          np.average(scores3)
```

Out[102]: 0.9409931719428926

```
In [103]: scores4 = cross_val_score(RandomForestClassifier(n_estimators=40), digit  
s.data, digits.target, cv=10)  
np.average(scores4)
```

Out[103]: 0.9443358162631904

Here we used `cross_val_score` to fine tune our random forest classifier and figured that having around 40 trees in random forest gives best result.

Exercise

Use iris flower dataset from sklearn library and use `cross_val_score` against following models to measure the performance of each. In the end figure out the model with best performance,

1. Logistic Regression
2. SVM
3. Decision Tree
4. Random Forest